

Generative Adversarial Networks

Deep Learning (DSE316/616)

Vinod K Kurmi
Assistant Professor, DSE

Indian Institute of Science Education and Research Bhopal

Oct 31, 2022



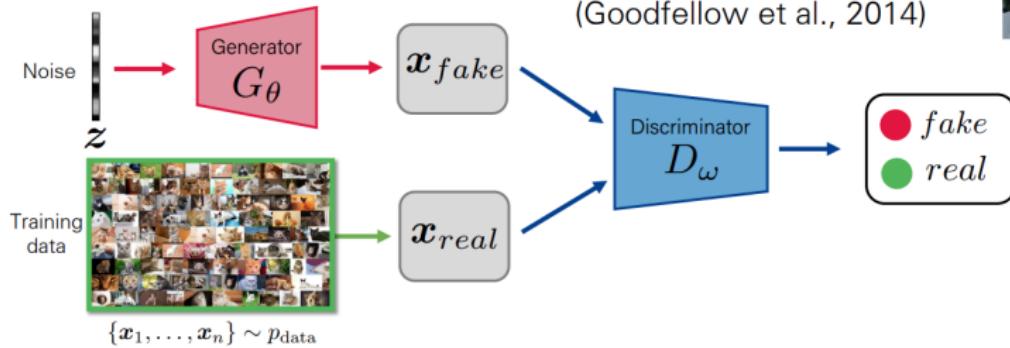
Disclaimer

- Much of the material and slides for this lecture were borrowed from
 - Bernhard Schölkopf's MLSS 2017 lecture,
 - Tommi Jaakkola's 6.867 class,
 - CMP784: Deep Learning Fall 2021 Erkut Erdem Hacettepe University
 - Fei-Fei Li, Andrej Karpathy and Justin Johnson's CS231n class
 - Hongsheng Li's ELEG5491 class
 - Tsz-Chiu Au slides
 - Mitesh Khapra Class notes

Genetive Adversarial Networks (GANs)

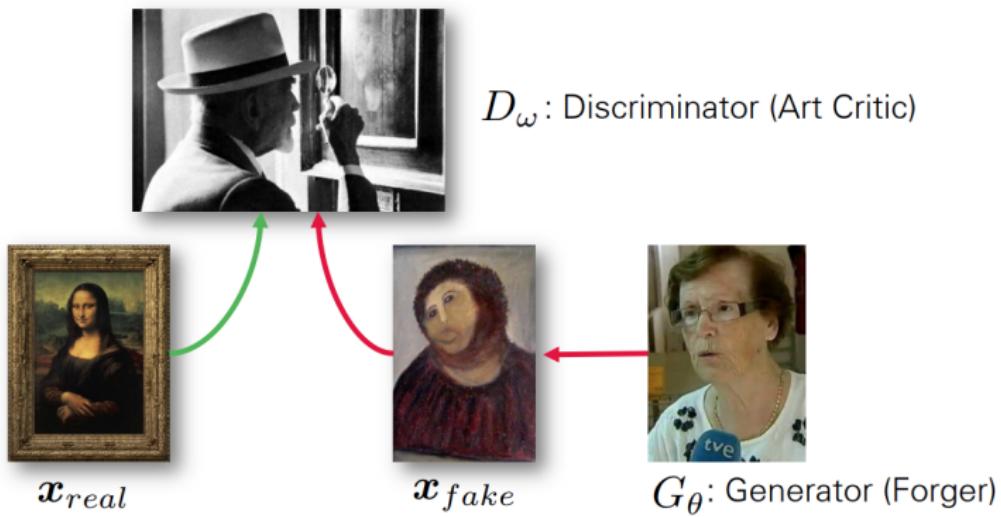
Genetive Adversarial Networks (GANs)

(Goodfellow et al., 2014)



- A game between a generator $G_\theta(\mathbf{z})$ and a discriminator $D_\omega(\mathbf{x})$
 - Generator tries to fool discriminator (i.e. generate realistic samples)
 - Discriminator tries to distinguish fake from real samples

Intuition behind GANs



GAN Training: Minimax Game

$$\min_{\theta} \max_{\omega} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D_{\omega}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}} [\log (1 - D_{\omega}(G_{\theta}(\mathbf{z})))]$$

Real data
↑

Noise vector used
to generate data
↑

Cross-entropy
loss for binary
classification

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z}} \log (1 - D(G(\mathbf{z})))$$

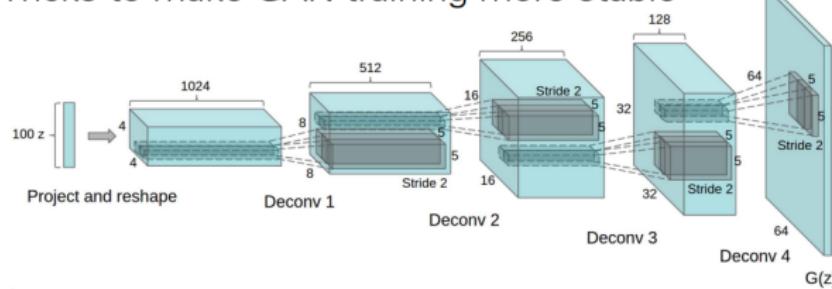
$$J^{(G)} = -\frac{1}{2} \mathbb{E}_{\mathbf{z}} \log D(G(\mathbf{z}))$$

Generator maximizes the log-probability
of the discriminator being mistaken

Deep Convolutional GANs (DCGAN)

(Radford et al., 2015)

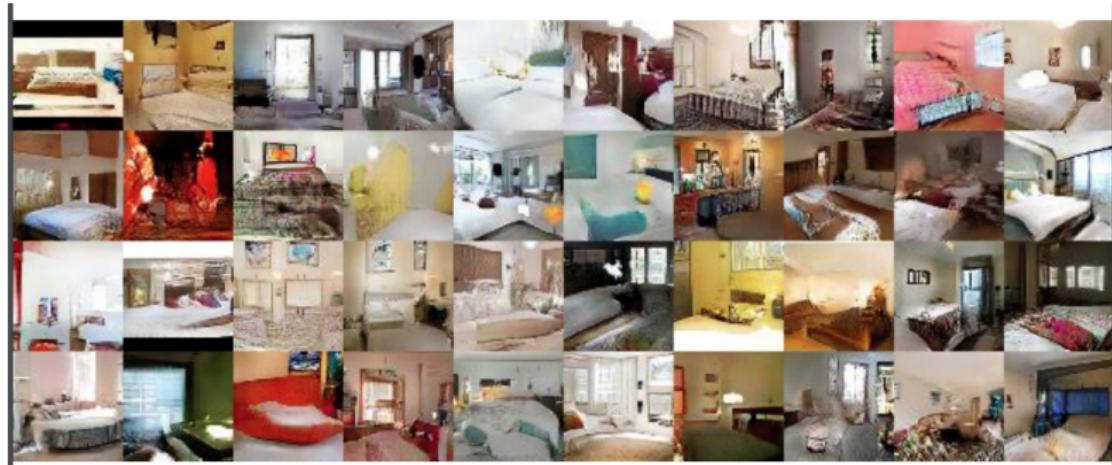
- Idea: Tricks to make GAN training more stable



- No fully connected layers
- Batch Normalization (Ioffe and Szegedy, 2015)
- Leaky Rectifier in D
- Use Adam (Kingma and Ba, 2015)
- Tweak Adam hyperparameters a bit ($\text{lr}=0.0002$, $b1=0.5$)

2/1

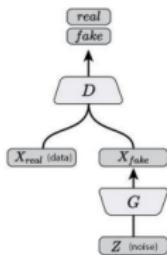
Deep Convolutional GANs (DCGAN)



Subclasses of GANs

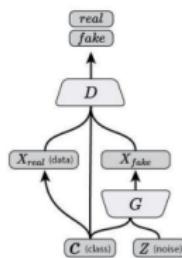
Vanilla GAN

Vanilla GAN
(Goodfellow, et al., 2014)

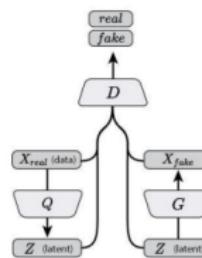


Discriminator Looks at Latent Variables

Conditional GAN
(Mirza & Osindero, 2014)

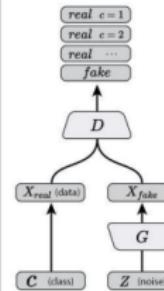


Bidirectional GAN
(Donahue, et al., 2016; Dumoulin, et al., 2016)

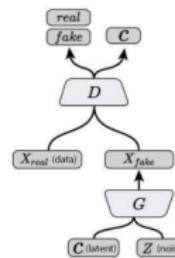


Discriminator Predicts Latent Variables

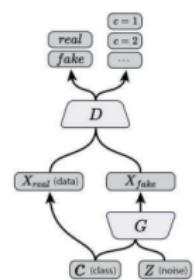
Semi-Supervised GAN
(Odena, 2016; Salimans, et al., 2016)



InfoGAN
(Chen, et al., 2016)

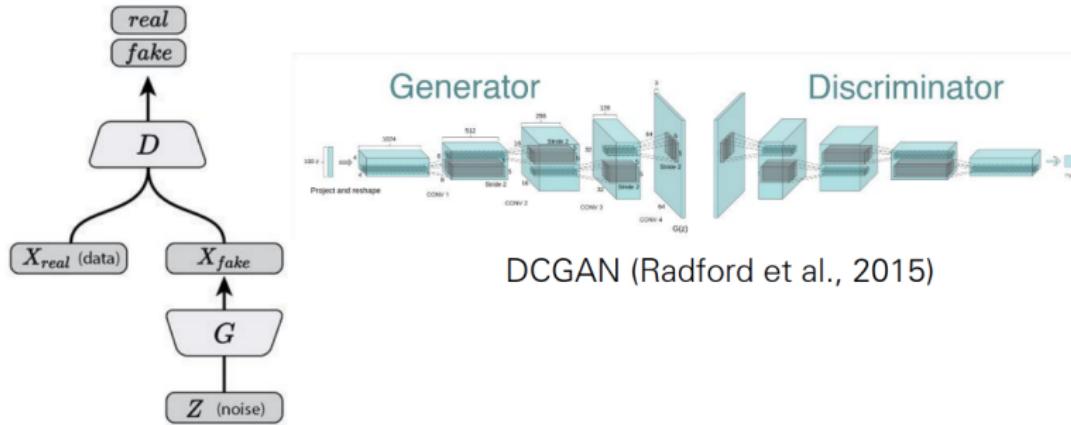


Auxiliary Classifier GAN
(Odena, et al., 2016)

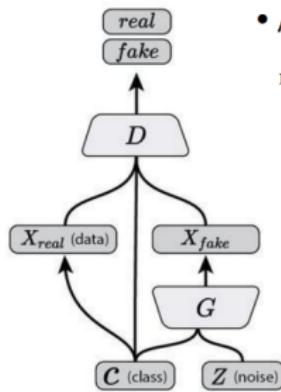


Vanilla GAN (Goodfellow et al., 2014)

- Updating the generator:

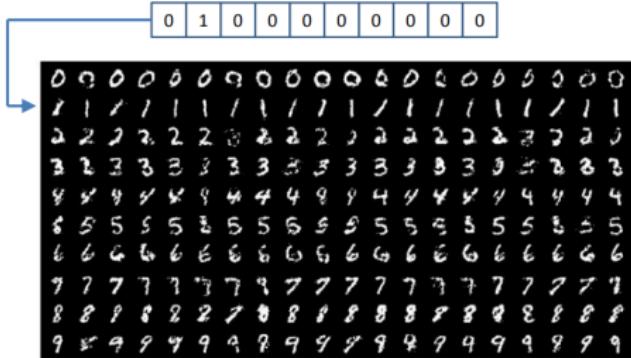


Conditional GAN (Mirza et al)

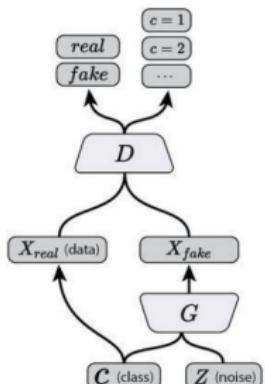


- Add conditional variables \mathbf{y} into G and D

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|\mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|\mathbf{y})))]$$



Auxiliary Classifier GAN(Odena et al)



- Every generated sample has a corresponding class label
$$L_S = E[\log P(S = real \mid X_{real})] + E[\log P(S = fake \mid X_{fake})]$$
$$L_C = E[\log P(C = c \mid X_{real})] + E[\log P(C = c \mid X_{fake})]$$
- D is trained to maximize $L_S + L_C$
- G is trained to maximize $L_C - L_S$
- Learns a representation for z that is independent of class label

Auxiliary Classifier GAN



monarch butterfly



goldfinch



daisy



redshank

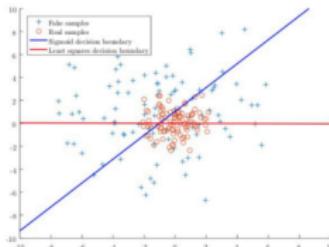


grey whale

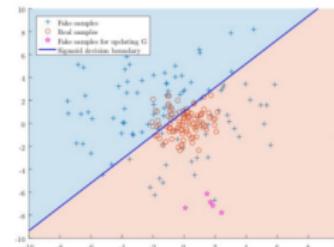
Least Squares GAN (LSGAN)

- Use a loss function that provides smooth and non-saturating gradient in discriminator D

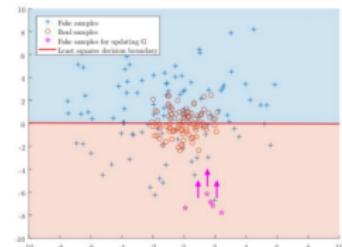
$$\min_D V_{\text{LSGAN}}(D) = \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [(D(\mathbf{x}) - b)^2] + \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [(D(G(\mathbf{z})) - a)^2]$$
$$\min_G V_{\text{LSGAN}}(G) = \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [(D(G(\mathbf{z})) - c)^2],$$



Decision boundaries of Sigmoid & Least Squares loss functions



Sigmoid decision boundary



Least Squares decision boundary

Least Squares GAN (LSGAN)



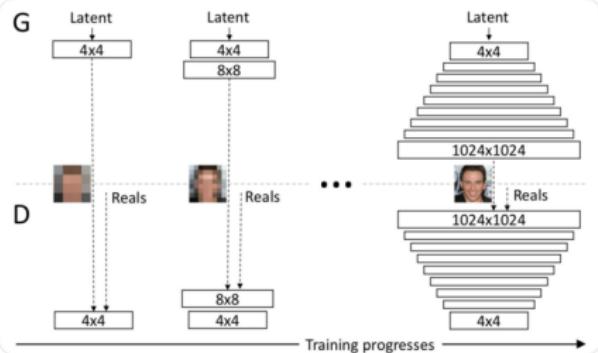
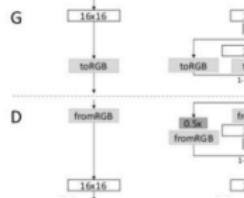
Church



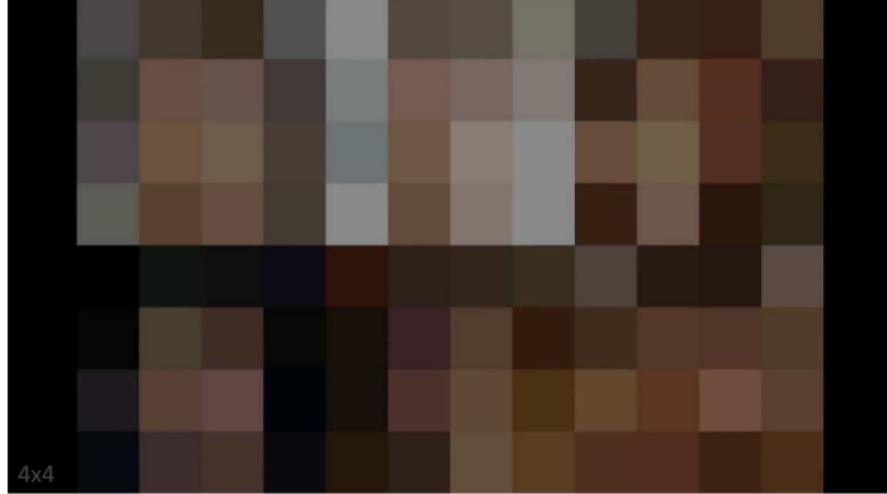
Kitchen

Progressive GANs(Karras et al., 2018)

- Progressively generate high-res images
- Multi-step training from low to high resolutions



Progressive GANs(Karras et al., 2018)



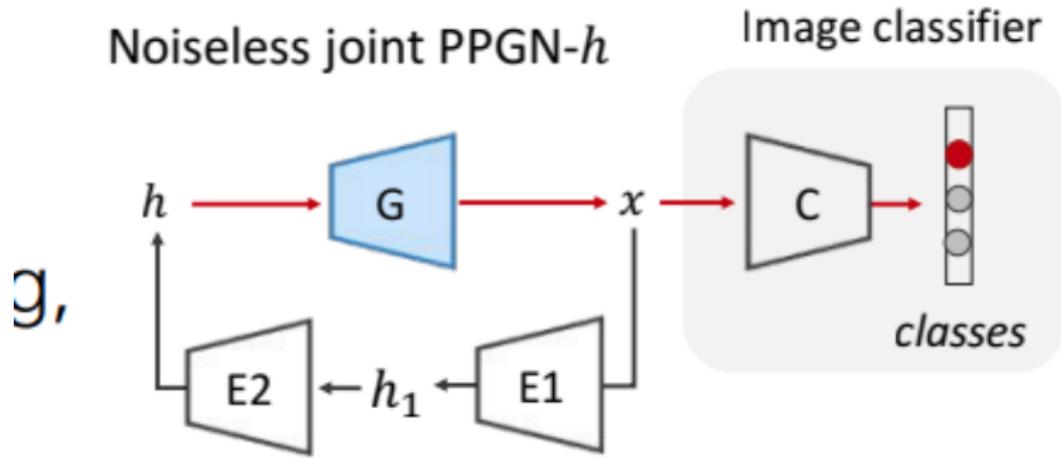
- Training process

Progressive GANs(Karras et al., 2018)

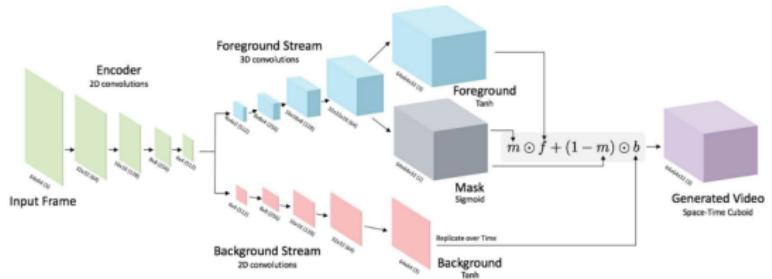


Class-specific Image Generation (Nguyen et al., 2016)

- Generates 227x227 realistic images from all ImageNet classes
- Combines adversarial training, moment matching, denoising autoencoders, and Langevin sampling



Video Generation (Vondrick et al., 2016)



Text-to-Image Synthesis (Zhang et al., 2016)

The small bird has a red head with feathers that fade from red to gray from head to tail



The petals of this flower are white with a large stigma



A unique yellow flower with no visible pistils protruding from the center



This flower is pink and yellow in color, with petals that are oddly shaped



This is a light colored flower with many different petals on a green stem



This flower is yellow and green in color, with petals that are ruffled



The flower have large petals that are pink with yellow on some of the petals



A flower that has white petals with some tones of yellow and green filaments



Text-to-Image Synthesis (Zhu et al., 2019)

This bird has a white throat and a dark yellow bill and grey wings.



This particular bird has a belly that is yellow and brown.



This bird is a lime green with greyish wings and long legs.



This yellow bird has a thin beak and jet black eyes and thin feet.



This bird has wings that are grey and has a white belly.



This bird has wings that are black and has a white belly.



This is a grey bird with a brown wing and a small orange beak.



This bird has a short brown bill, a white eyering, and a medium brown crown.



Single Image Super-Resolution ((Ledig et al., 2016))

- Combine content loss with adversarial loss

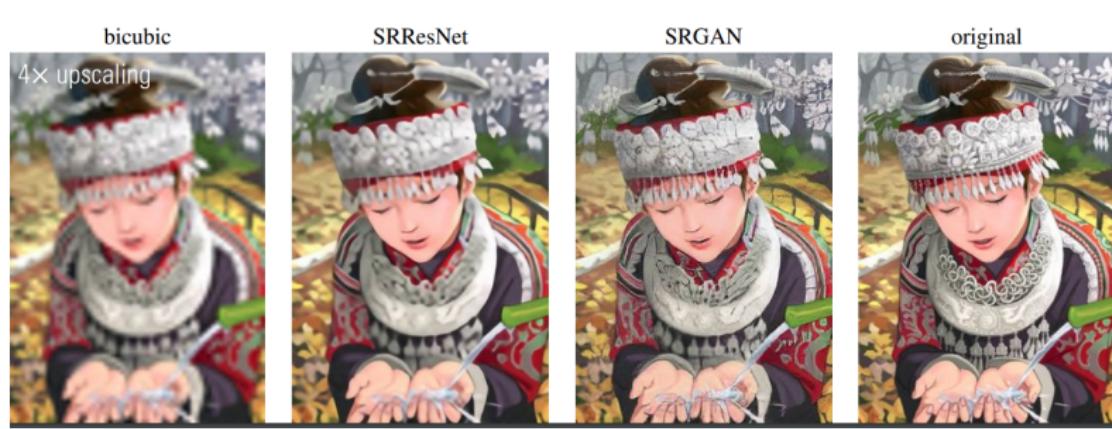


Image Inpainting (Pathak et al., 2016)

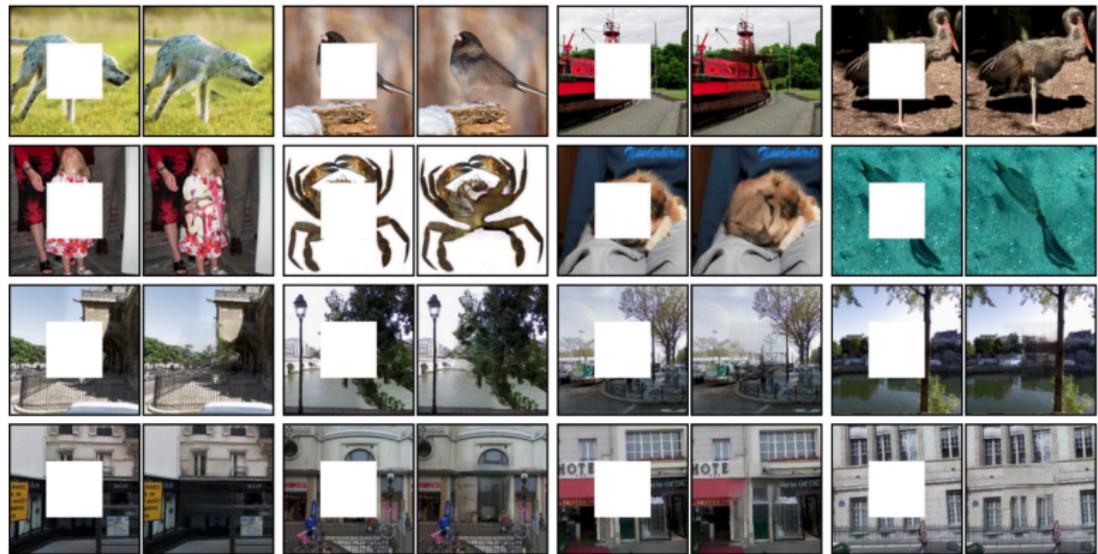


Image to Image Translation (Pix2Pix)



(Isola et al. 2016)



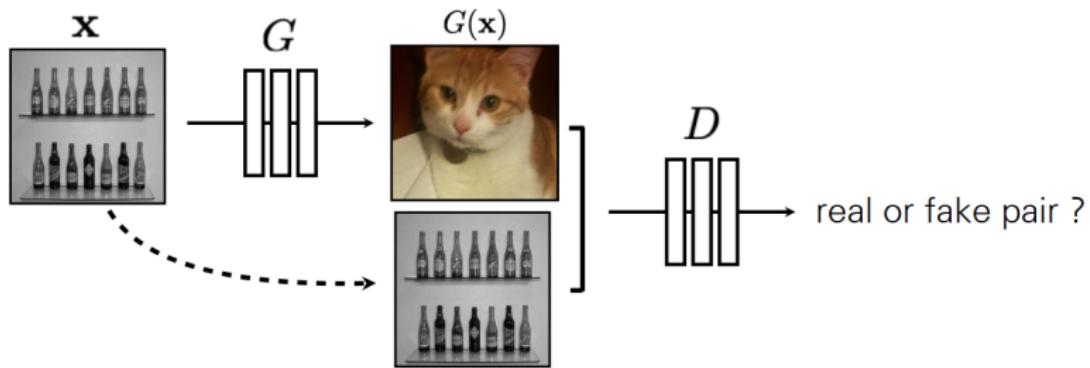
Image to Image Translation (Pix2Pix)



(Isola et al. 2016)



Image to Image Translation (Pix2Pix)



$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$

Image to Image Translation (Pix2Pix)

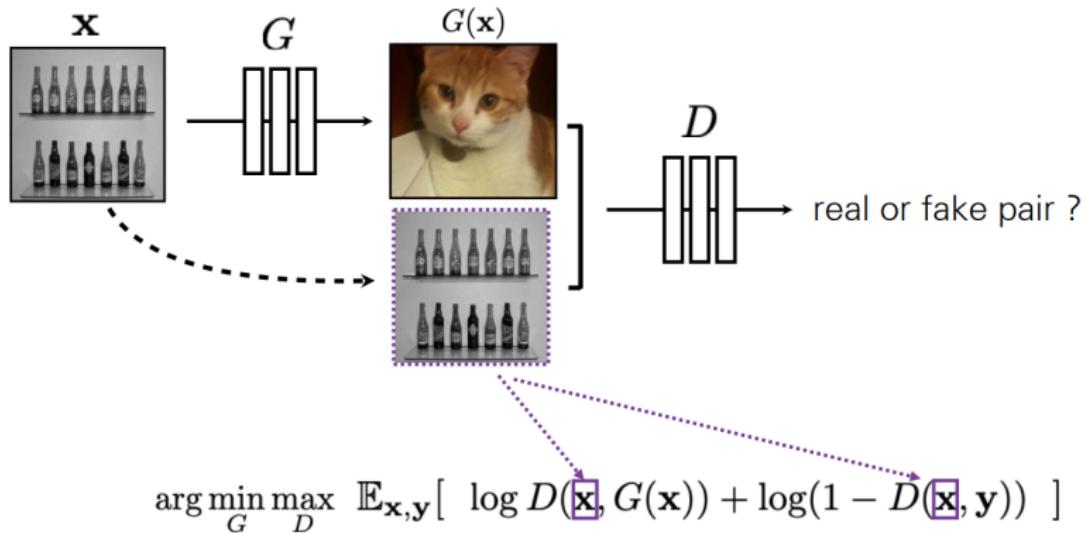


Image to Image Translation (Pix2Pix)

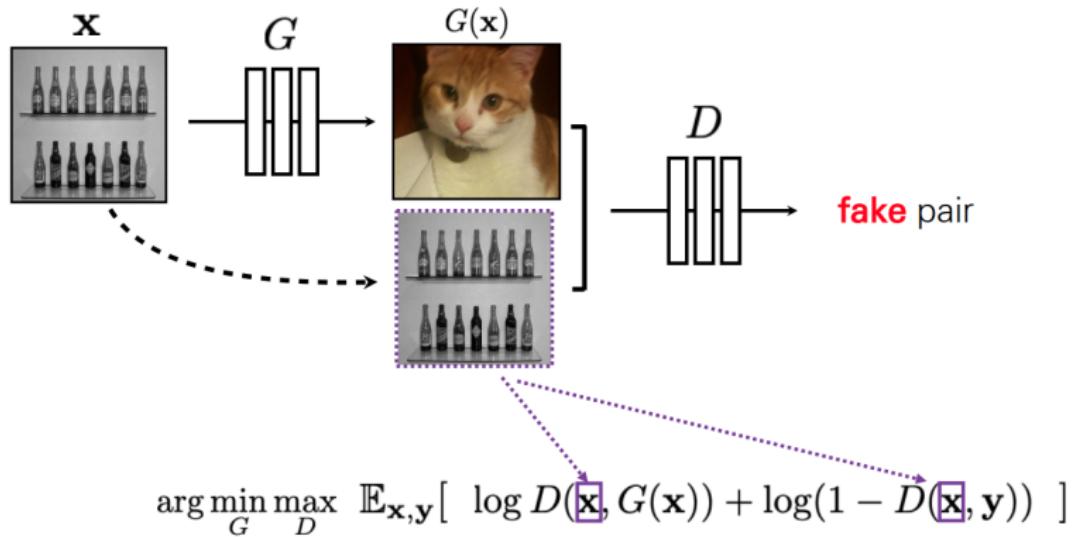


Image to Image Translation (Pix2Pix)

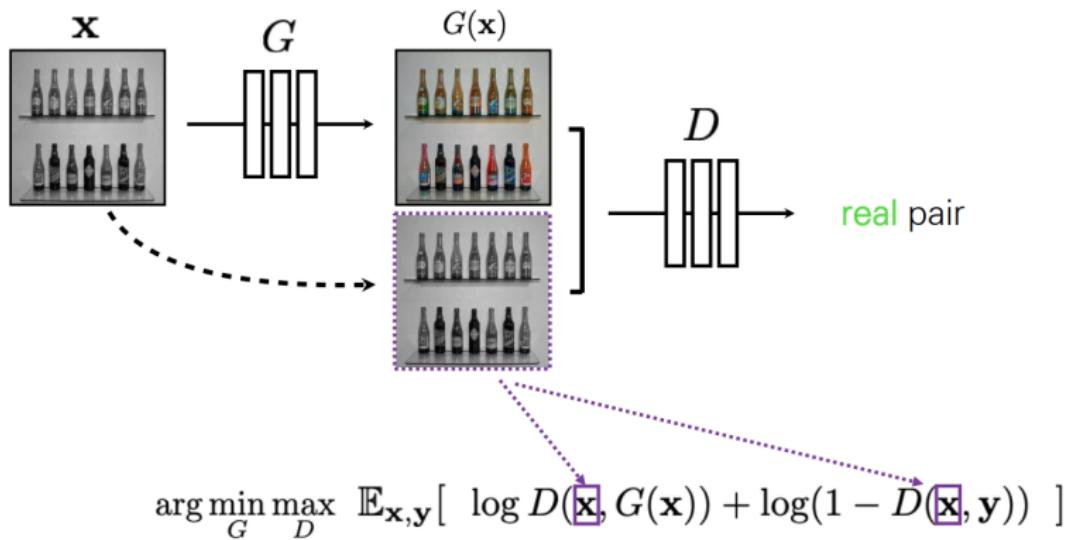
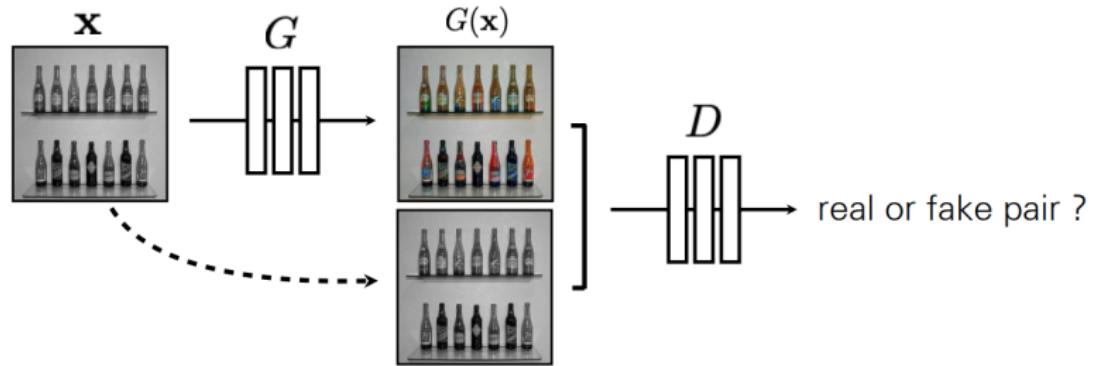
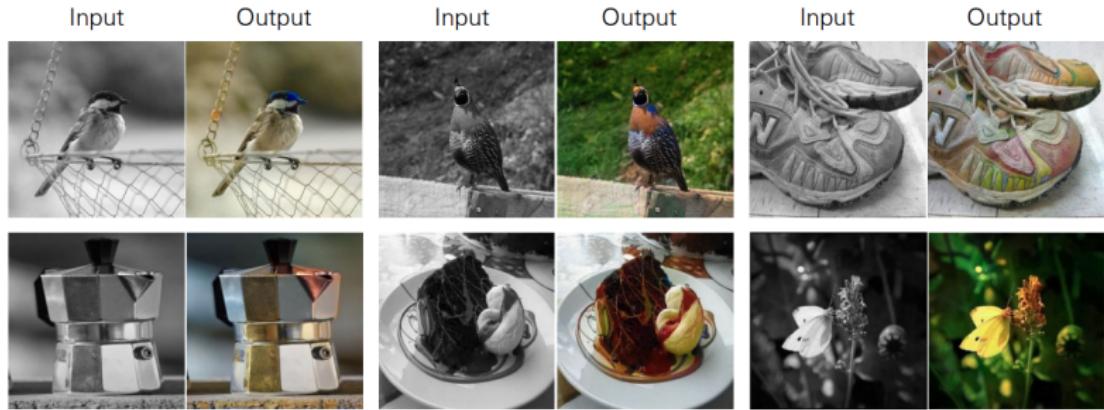


Image to Image Translation (Pix2Pix)

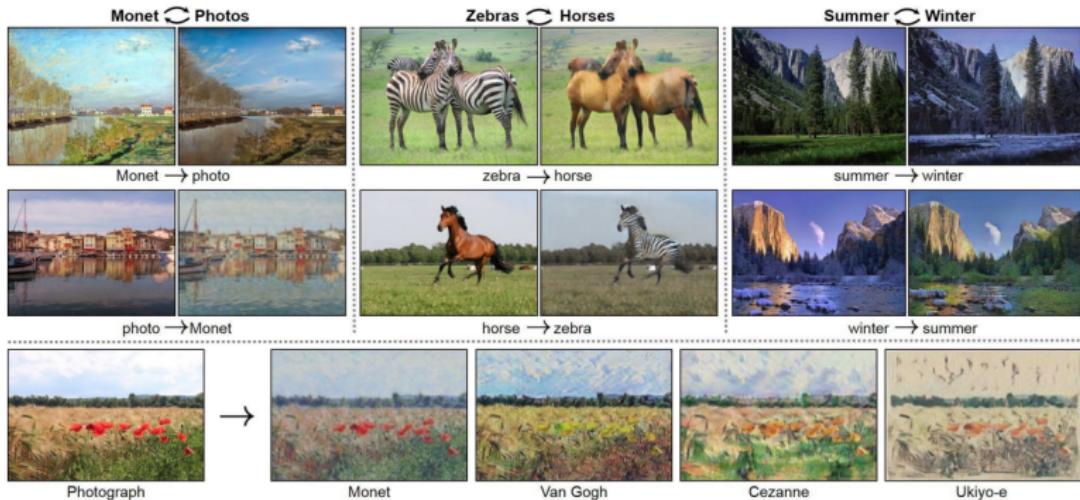


$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))]$$

BW → Color



CycleGAN: Pix2Pix w/o input-output pairs



(Zhu et al. 2017)

Paired data

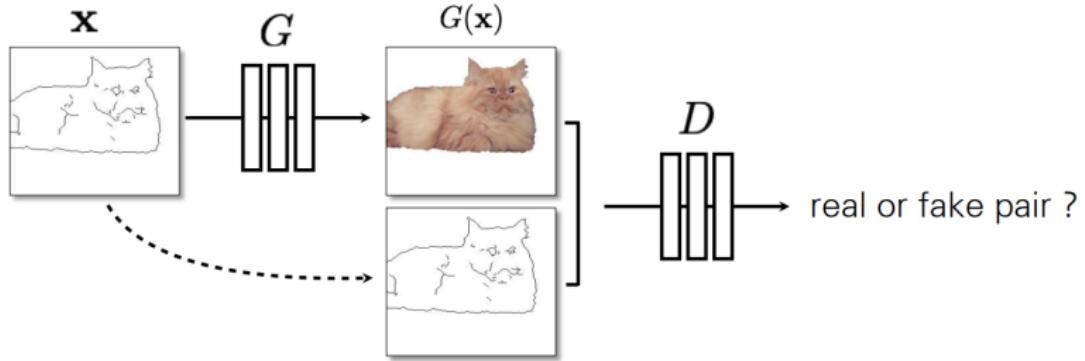
x_i	y_i
{  ,  }	
{  ,  }	
{  ,  }	
⋮	

Paired data

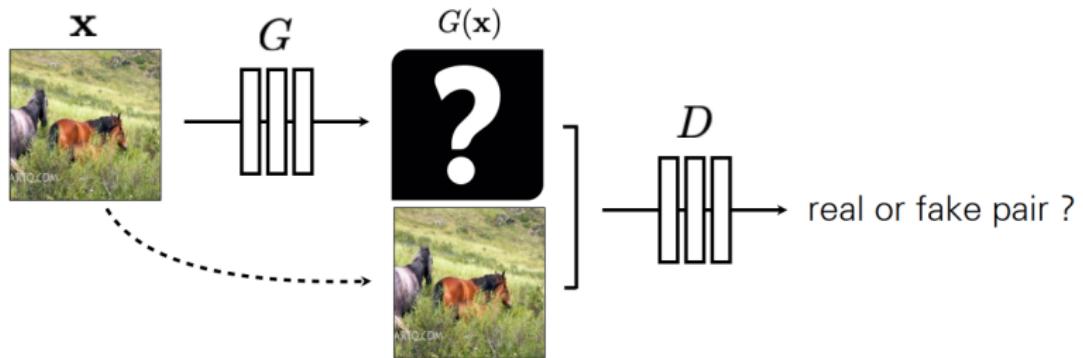
x_i	y_i
{  , }	{  }
{  , }	{  }
{  , }	{  }
⋮	

Unpaired data

X	Y
{  }	{  }
{  }	{  }
{  }	{  }
⋮	⋮

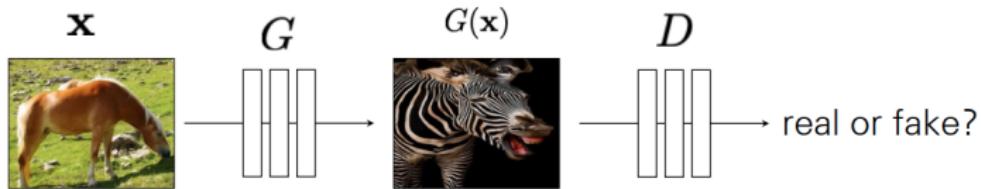


$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))]$$



$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))]$$

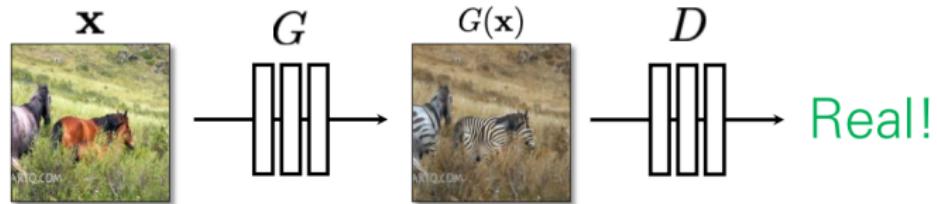
No input-output pairs!

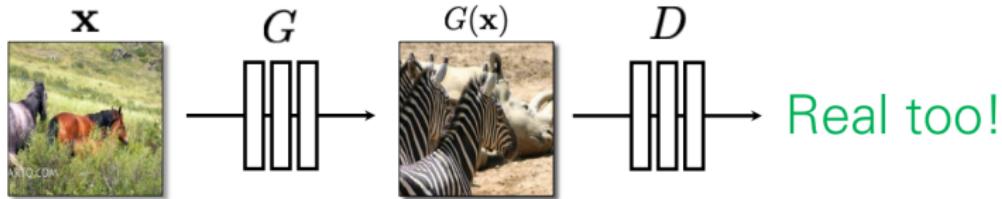


$$\arg \min_G \max_D \mathbb{E}_{\mathbf{x}, \mathbf{y}} [\log D(G(\mathbf{x})) + \log(1 - D(\mathbf{y}))]$$

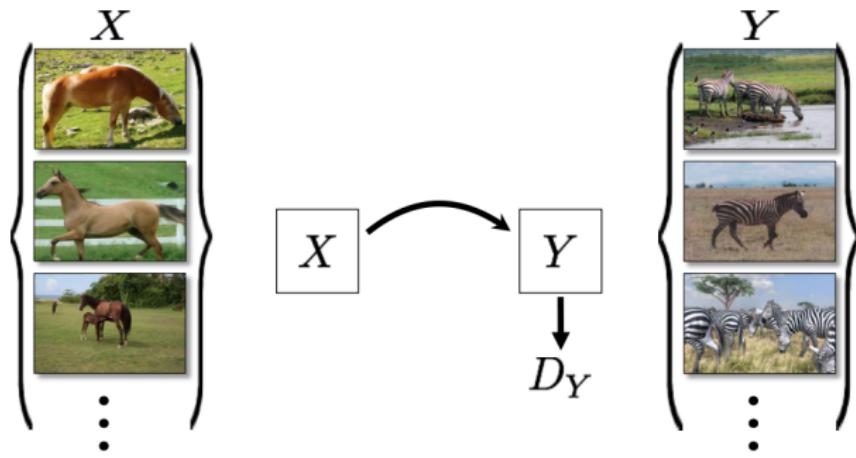
Usually loss functions check if output matches a target instance

GAN loss checks if output is part of an admissible set

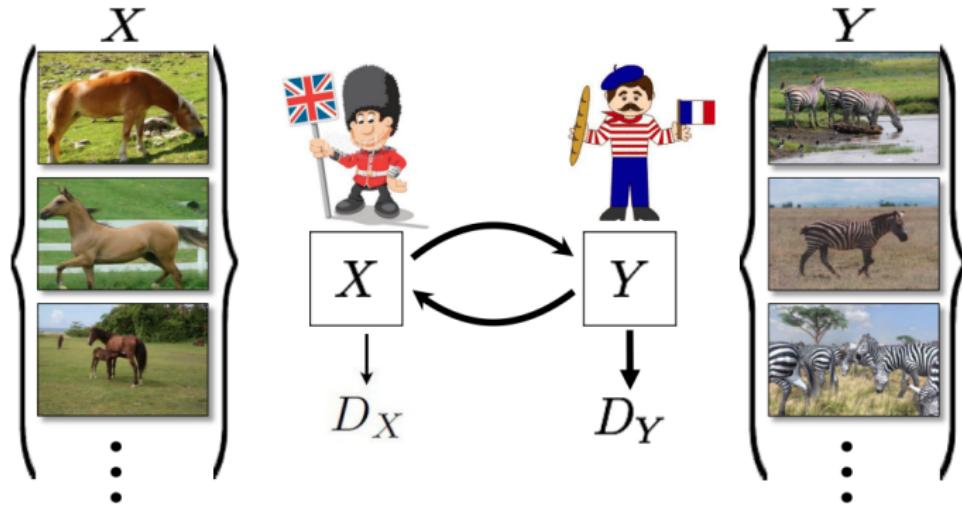




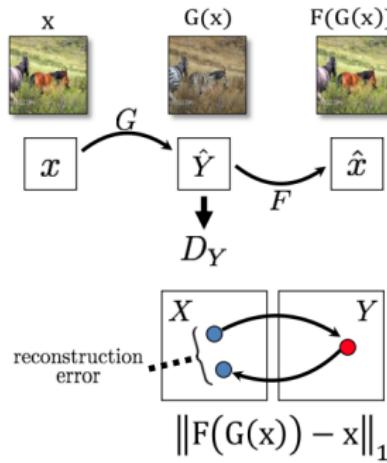
Nothing to force output to correspond to input



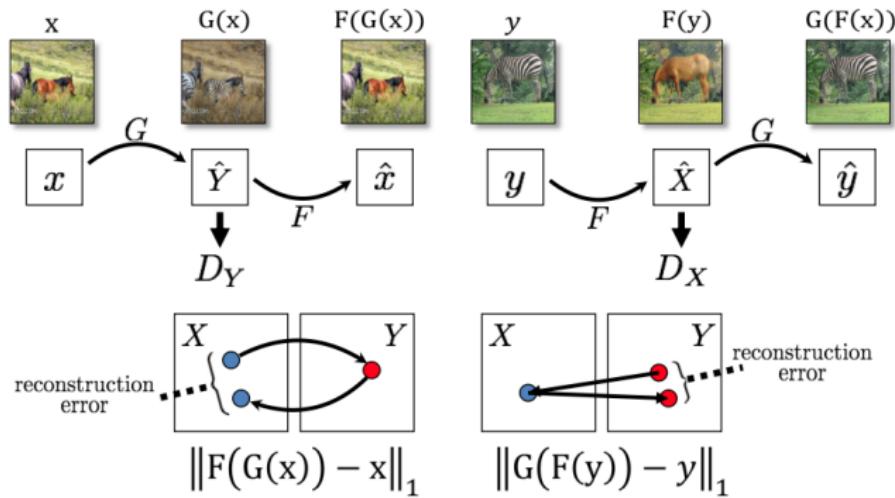
[Zhu et al. 2017], [Yi et al. 2017], [Kim et al. 2017]



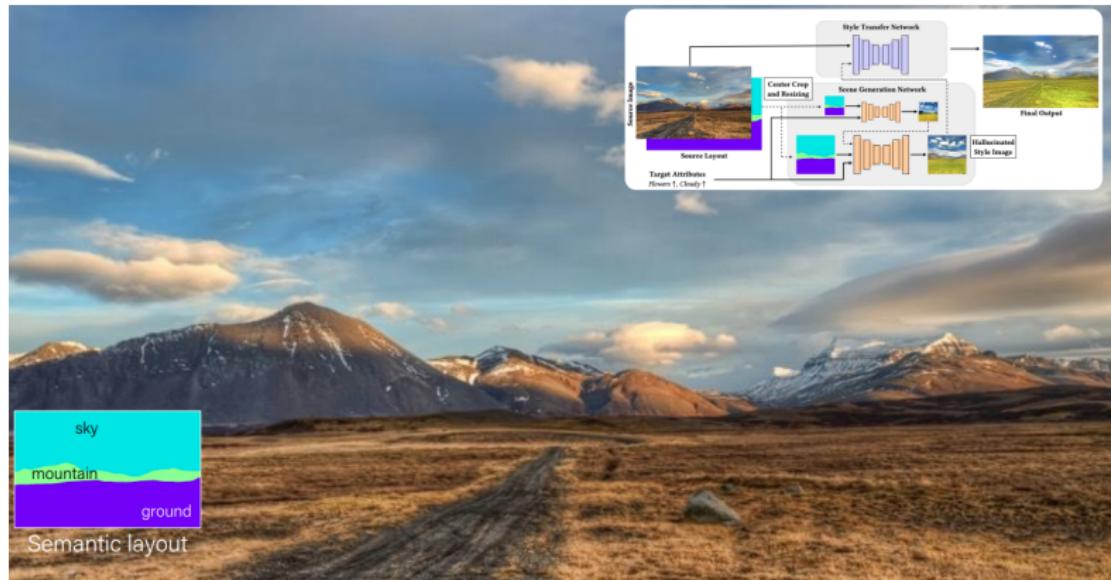
Cycle-Consistent Adversarial Networks



Cycle-Consistent Adversarial Networks



Cycle Consistency Loss



Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]

Cycle Consistency Loss



Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]

Semantic Image Synthesis (SPADE)



Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]



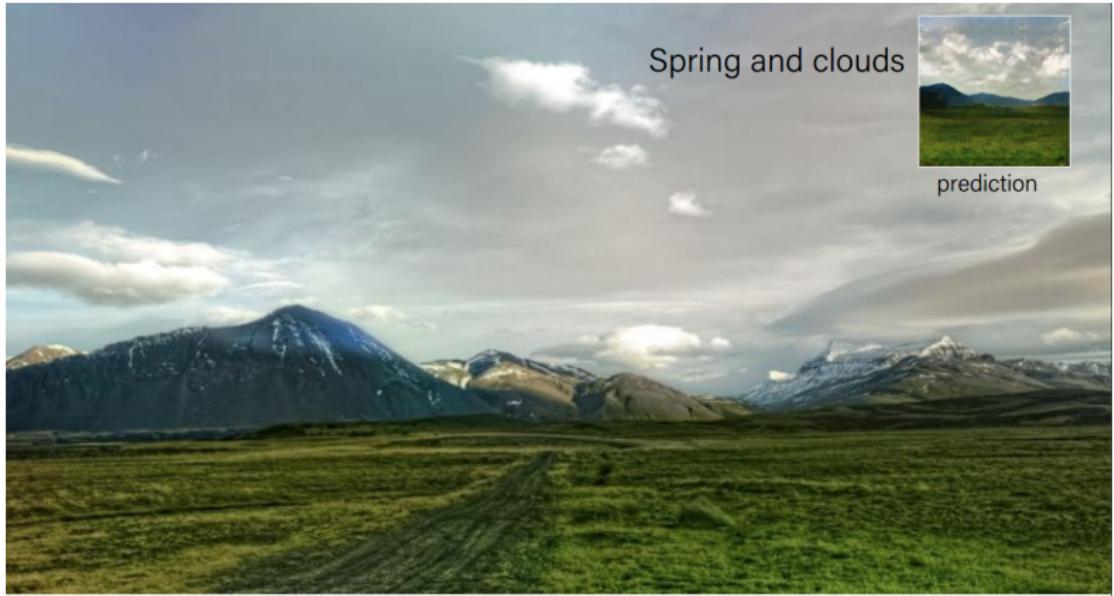
Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]



Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]



Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]



Spring and clouds



prediction

Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]



Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2020]



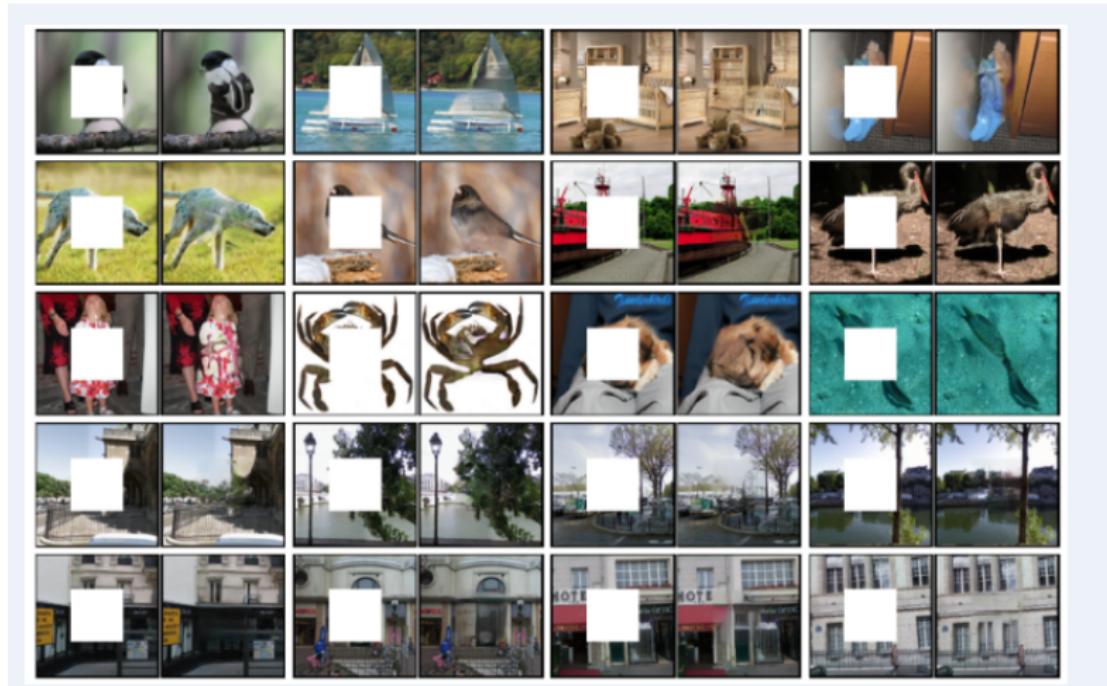
Manipulating Attributes of Natural Scenes via Hallucination [Karacan et al., 2018]

Context Encoder: Feature Learning by Inpainting (Pathak et al)

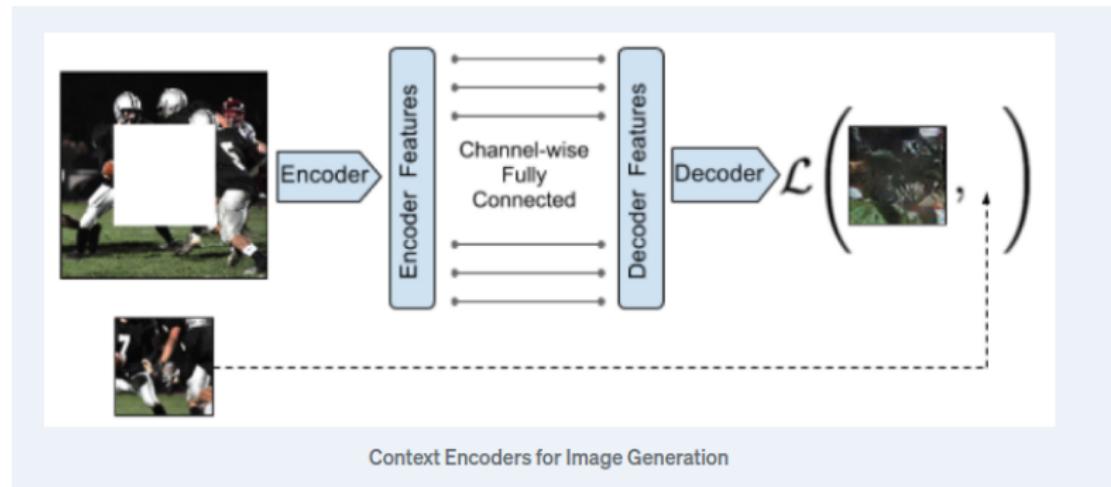
Context Encoder: Feature Learning by Inpainting (Pathak et al)

- A CNN, called Context Encoders, is trained to generate the contents of an arbitrary image region conditioned on its surroundings.
- Both a standard pixel-wise reconstruction loss, as well as a reconstruction plus an adversarial loss, are tested.
- The Context Encoders can be used for feature learning. The learnt features can be used pre-training on classification, detection, and segmentation tasks, which is a kind of self-supervised learning.

Context Encoder: Feature Learning by Inpainting



Context Encoder: Feature Learning by Inpainting



Context Encoder: Feature Learning by Inpainting

- For each ground truth image x , the proposed context encoder F produces an output $F(x)$.
- Let \hat{M} be a binary mask corresponding to the dropped image region with a value of 1 wherever a pixel was dropped and 0 for input pixels.
- The reconstruction loss is a normalized masked L2 distance:

$$\mathcal{L}_{rec}(x) = \|\hat{M} \odot (x - F((1 - \hat{M}) \odot x))\|_2^2,$$

- where \odot is the element-wise product operation.
- L1 and L2 losses have no significant difference, it often fails to capture any high frequency details, often prefer a blurry solution, over highly accurate textures.

Context Encoder: Feature Learning by Inpainting

- The adversarial loss tries to make prediction look real, and has the effect of picking a particular mode from the distribution.
- Only the generator (not the discriminator) is conditioned on context when trained using GAN. The adversarial loss for context encoders, L_{adv} , is:

$$\begin{aligned}\mathcal{L}_{adv} = \max_D \quad & \mathbb{E}_{x \in \mathcal{X}} [\log(D(x)) \\ & + \log(1 - D(F((1 - \hat{M}) \odot x)))],\end{aligned}$$

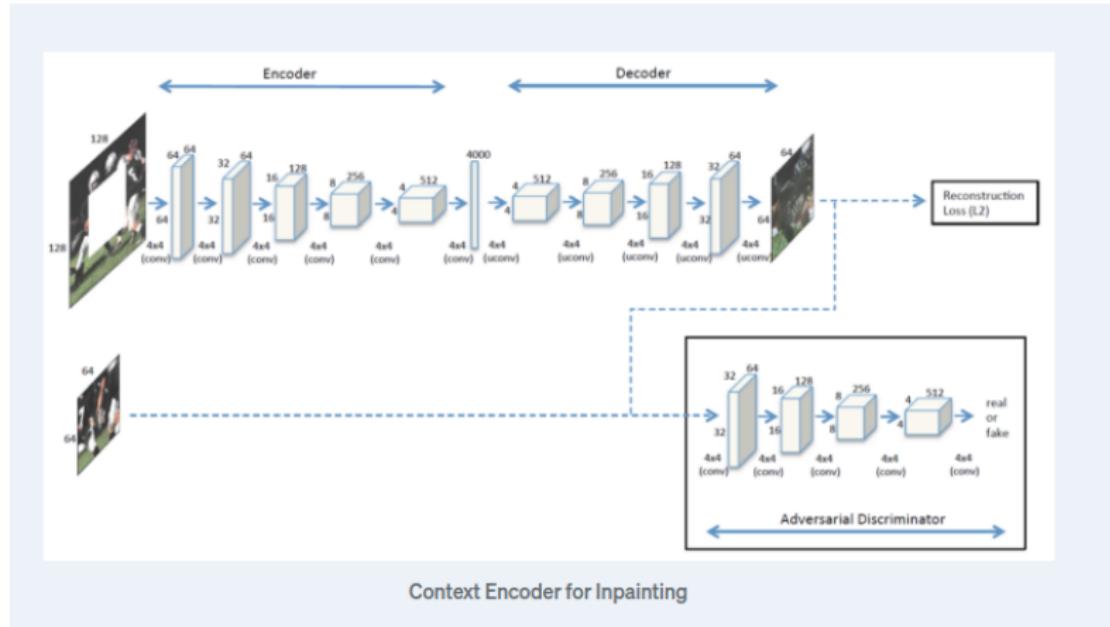
- Both F and D are optimized jointly using alternating SGD.

Context Encoder: Feature Learning by Inpainting

- The overall loss function is:

$$\mathcal{L} = \lambda_{rec}\mathcal{L}_{rec} + \lambda_{adv}\mathcal{L}_{adv}.$$

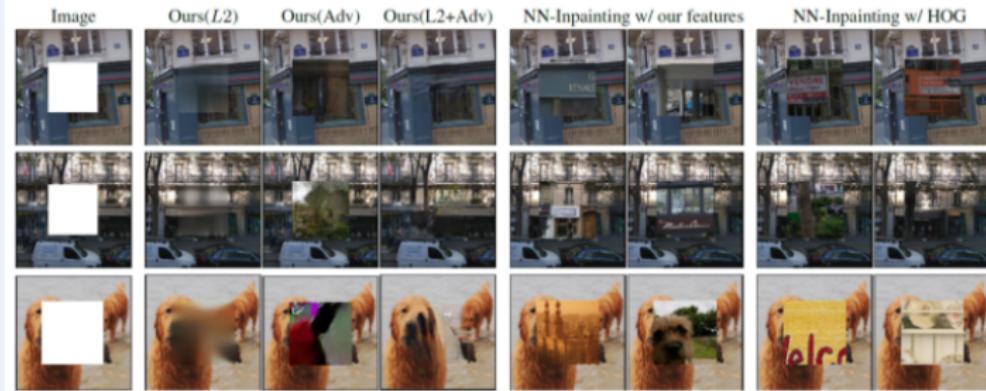
Context Encoder: Feature Learning by Inpainting



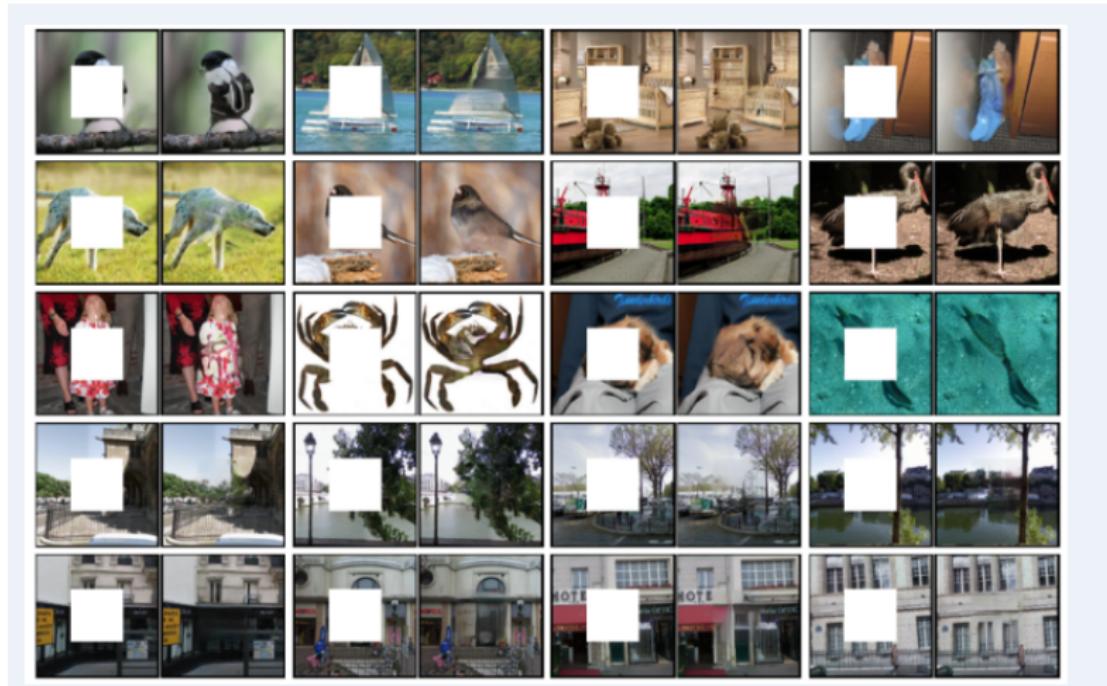
Context Encoder: Feature Learning by Inpainting



Context Encoder: Feature Learning by Inpainting



Context Encoder: Feature Learning by Inpainting



Applications of Adversarial loss

- Image generation: colonization, super resolution.....
- Domain Adaptation
- Fair learning
- Active learning
-

Next Class

- Self supervised learning
- Zero shot learning
- One shot learning
- Few shot learning
- Domain adaptation
- Active learning
- Bayesian learning
- Fair learning